

# Introduction

[www.it-ebooks.info](http://www.it-ebooks.info) . NLTK is often used for rapid prototyping of text processing programs and can even be used in production applications.

## Tokenizing text into sentences

Instructions on how to do so are available at <http://nltk.org/install.html>. Make sure that `tokenizers/punkt.zip` is in the data directory and has been unpacked so that there's a file at `tokenizers/punkt/PY3/english.pickle`.

## Tokenizing sentences into words

These differ from `TreebankWordTokenizer` by how they handle punctuation and contractions, but they all inherit from `TokenizerI`. 11 [www.it-ebooks.info](http://www.it-ebooks.info) .

## Tokenizing sentences using regular expressions

As regular expressions can get complicated very quickly, I only recommend using them if the word tokenizers covered in the previous recipe are unacceptable. We'll create an instance of `RegexpTokenizer`, giving it a regular expression string to use for matching tokens: 

```
>>> from nltk.tokenize import RegexpTokenizer
>>> tokenizer = RegexpTokenizer("[\w']+")
>>> tokenizer.tokenize("Can't is a contraction.")
```

`['Can't', 'is', 'a', 'contraction']` There's also a simple helper function you can use if you don't want to instantiate the class, as shown in the following code: 

```
>>> from nltk.tokenize import regexp_tokenize
>>> regexp_tokenize("Can't is a contraction.", "[\w']+")
```

`['Can't', 'is', 'a', 'contraction']` Now we finally have something that can treat contractions as whole words, instead of splitting them into tokens.

## Training a sentence tokenizer

As you can see, this isn't your standard paragraph of sentences formatting, which makes it a perfect case for training a sentence tokenizer. This difference is a good demonstration of why it can be useful to train your own sentence tokenizer, especially when your text isn't in the typical paragraph-sentence structure.

## Filtering stopwords in a tokenized sentence

For example: `>>> stopwords.words('dutch') ['de', 'en', 'van', 'ik', 'te', 'dat', 'die', 'in', 'een', 'hij', 'het', 'niet', 'zijn', 'is', 'was', 'op', 'aan', 'met', 'als', 'voor', 'had', 'er', 'maar', 'om', 'hem', 'dan', 'zou', 'of', 'wat', 'mijn', 'men', 'dit', 'zo', 'door', 'over', 'ze', 'zich', 'bij', 'ook', 'tot', 'je', 'mij', 'uit', 'der', 'daar', 'haar', 'naar', 'heb', 'hoe', 'heeft', 'hebben', 'deze', 'u', 'want', 'nog', 'zal', 'me', 'zij', 'nu', 'ge', 'geen', 'omdat', 'iets', 'worden', 'toch', 'al', 'waren', 'veel', 'meer', 'doen', 'toen', 'moet', 'ben', 'zonder', 'kan', 'hun', 'dus', 'alles', 'onder', 'ja', 'eens', 'hier', 'wie', 'werd', 'altijd', 'doch', 'wordt', 'wezen', 'kunnen', 'ons', 'zelf', 'tegen', 'na', 'reeds', 'wil', 'kon', 'niets', 'uw', 'iemand', 'geweest', 'andere']` 17 [www.it-ebooks.info](http://www.it-ebooks.info) . As such, it has a `words()` method that can take a single argument for the file ID, which in this case is 'english', referring to a file containing a list of English stopwords.

## Looking up Synsets for a word in WordNet

You can trace the entire path from entity down to cookbook using the `hypernym_paths()` method, as follows: `>>> syn.hypernym_paths() [[Synset('entity.n.01'), Synset('physical_entity.n.01'), Synset('object.n.01'), Synset('whole.n.02'), Synset('artifact.n.01'), Synset('creation.n.02'), Synset('product.n.02'), Synset('work.n.02'), Synset('publication.n.01'), Synset('book.n.01'), Synset('reference_book.n.01'), Synset('cookbook.n.01')]]` 19 [www.it-ebooks.info](http://www.it-ebooks.info) . The `name()` method will give you a unique name for the Synset, which you can use to get the Synset directly: `>>> wordnet.synset('cookbook.n.01')` `Synset('cookbook.n.01')` The `definition()` method should be self-explanatory.

## Looking up lemmas and synonyms in WordNet

`synonyms.append(lemma.name())` `>>> len(synonyms)` 38 21 [www.it-ebooks.info](http://www.it-ebooks.info) Tokenizing Text and WordNet Basics As you can see, there appears to be 38 possible synonyms for the word 'book'. for lemma in `syn.lemmas()`: ...

## Calculating WordNet Synset similarity

This tree can be used for reasoning about the similarity between the Synsets it contains. The closer the two Synsets are in the tree, the more similar they are.

## Discovering word collocations

f This can be used to inform your choice for min\_score. Discovering collocations in this list of words means that we'll find common phrases that occur frequently throughout the text.

## Introduction

All of these methods can be very useful for preprocessing text before search indexing, document classification, and text analysis. [www.it-ebooks.info](http://www.it-ebooks.info) .

## Stemming words

For example, the stem of cooking is cook, and a good stemming algorithm knows that the ing suffix can be removed. The resulting stem is often a shorter word, or at least a common form of the word, which has the same root meaning.

## Lemmatizing words with WordNet

So unlike stemming, you are always left with a valid word that means the same thing. However, the word you end up with can be completely different.

## Replacing words matching regular expressions

The following code can be found in the `replacers.py` module in the book's code bundle and is meant to be imported, not typed into the console:

```
import re
replacement_patterns = [ (r'won't', 'will not'),
(r'can't', 'cannot'), (r'i'm', 'i am'), (r'ain't', 'is not'), (r'(\w+)\ll', '\g<1> will'), (r'(\w+)n't', '\g<1> not'),
(r'(\w+)\ve', '\g<1> have'), (r'(\w+)\s', '\g<1> is'), (r'(\w+)\re', '\g<1> are'), (r'(\w+)\d', '\g<1> would') ]
class RegexpReplacer(object):
    def __init__(self, patterns=replacement_patterns):
        self.patterns = [(re.compile(regex), repl) for (regex, repl) in patterns]
    def replace(self, text):
        s = text
        for (pattern, repl) in self.patterns:
            s = re.sub(pattern, repl, s)
        return s
```

35 [www.it-ebooks.info](http://www.it-ebooks.info)

Replacing and Correcting Words How it works... This will be a list of tuple pairs, where the first element is the pattern to match with and the second element is the replacement.

## Removing repeating characters

This code can be found in `replacers.py` in the book's code bundle and is meant to be imported:

```
import re
class RepeatReplacer(object):
    def __init__(self):
        self.repeat_regexp = re.compile(r'(\w*)(\w)\2(\w*)')
        self.repl = r'\1\2\3'
    def replace(self, word):
        repl_word = self.repeat_regexp.sub(self.repl, word)
        if repl_word != word:
            return self.replace(repl_word)
        else:
            return repl_word
```

37 [www.it-ebooks.info](http://www.it-ebooks.info) Replacing and Correcting Words

And now some example use cases:

```
>>> from replacers import RepeatReplacer
>>> replacer = RepeatReplacer()
>>> replacer.replace('loooooove')
'love'
>>> replacer.replace('oooooh')
'oh'
>>> replacer.replace('goose')
'gose'
```

How it works... Here is the WordNet-augmented version:

```
import re
from nltk.corpus import wordnet
class RepeatReplacer(object):
    def __init__(self):
        self.repeat_regexp = re.compile(r'(\w*)(\w)\2(\w*)')
        self.repl = r'\1\2\3'
```

38 [www.it-ebooks.info](http://www.it-ebooks.info) .

## Spelling correction with Enchant

If not, we will look up the suggested alternatives and return the best match using `nltk.metrics.edit_distance()`:

```
import enchant
from nltk.metrics import edit_distance
class SpellingReplacer(object):
    def __init__(self, dict_name='en', max_dist=2):
        self.spell_dict = enchant.Dict(dict_name)
        self.max_dist = max_dist
    def replace(self, word):
        if self.spell_dict.check(word):
            return word
        suggestions = self.spell_dict.suggest(word)
        if suggestions and edit_distance(word, suggestions[0]) <= self.max_dist:
            return suggestions[0]
        else:
            return word
```

The preceding class can be used to correct English spellings, as follows:

```
>>> from replacers import SpellingReplacer
>>> replacer = SpellingReplacer()
>>> replacer.replace('cookbok')
'cookbook'
```

How it works... You could then create a dictionary augmented with your personal word list as follows:

```
>>> d = enchant.Dict('en_US')
>>> d.check('nltk')
False
>>> d = enchant.DictWithPWL('en_US', 'mywords.txt')
>>> d.check('nltk')
True
```

To use an augmented dictionary with our `SpellingReplacer` class, we can create a subclass in `replacers.py` that takes an existing spelling dictionary:

```
class CustomSpellingReplacer(SpellingReplacer):
    def __init__(self, spell_dict, max_dist=2):
        self.spell_dict = spell_dict
        self.max_dist = max_dist
```

This `CustomSpellingReplacer` class will not replace any words that you put into `mywords.txt`:

```
>>> from replacers import CustomSpellingReplacer
>>> d = enchant.DictWithPWL('en_US', 'mywords.txt')
>>> replacer = CustomSpellingReplacer(d)
>>> replacer.replace('nltk')
'nltk'
```

See also The previous recipe covered an extreme form of spelling correction by replacing repeating characters.

## Replacing synonyms

CSV synonym replacement The CswWordReplacer class extends WordReplacer in replacers.py in order to construct the word\_map dictionary from a CSV file: import csv class CswWordReplacer(WordReplacer): def \_\_init\_\_(self, fname): word\_map = {} for line in csv.reader(open(fname)): word, syn = line word\_map[word] = syn super(CswWordReplacer, self).\_\_init\_\_(word\_map) Your CSV file should consist of two columns, where the first column is the word and the second column is the synonym meant to replace it. We'll first create a WordReplacer class in replacers.py that takes a word replacement mapping: class WordReplacer(object): def \_\_init\_\_(self, word\_map): self.word\_map = word\_map def replace(self, word): return self.word\_map.get(word, word) Then, we can demonstrate its usage for simple word replacement: >>> from replacers import WordReplacer >>> replacer = WordReplacer({'bday': 'birthday'}) >>> replacer.replace('bday') 'birthday' >>> replacer.replace('happy') 'happy' How it works...

## Replacing negations with antonyms

To do this, we will create an AntonymReplacer class in replacers.py as follows: from nltk.corpus import wordnet class AntonymReplacer(object): def replace(self, word, pos=None): antonyms = set() for syn in wordnet.synsets(word, pos=pos): for lemma in syn.lemmas(): for antonym in lemma.antonyms(): antonyms.add(antonym.name()) if len(antonyms) == 1: return antonyms.pop() else: return None def replace\_negations(self, sent): i, l = 0, len(sent) words = [] while i < l: word = sent[i] if word == 'not' and i+1 < l: ant = self.replace(sent[i+1]) if ant: words.append(ant) i += 2 continue words.append(word) i += 1 return words 46 www.it-ebooks.info Chapter 2 Now, we can tokenize the original sentence into ["let's", 'not', 'uglify', 'our', 'code'] and pass this to the replace\_negations() function. If not is found, then we try to find an antonym for the next word using replace().

## Introduction

This chapter will introduce you to many more corpora. This information is essential for future chapters when we'll need to access the corpora as training data.

## Setting up a custom corpus

Create this corpora directory within the nltk\_data directory, so that the path is ~/nltk\_data/corpora. For simplicity, I'll refer to the directory as ~/nltk\_data.

## Creating a wordlist corpus

In the Filtering stopwords in a tokenized sentence recipe in Chapter 1, Tokenizing Text and WordNet Basics, we saw that it had one wordlist file for each language, and you could access the words for that language by calling `stopwords.words(fileid)`. There's one file with 850 basic words, and another list with over 200,000 known English words, as shown in the following code: `>>> from nltk.corpus import words >>> words.fileids() ['en', 'en-basic'] >>> len(words.words('en-basic')) 850 >>> len(words.words('en')) 234936` See also The Filtering stopwords in a tokenized sentence recipe in Chapter 1, Tokenizing Text and WordNet Basics, has more details on using the stopwords corpus.

## Creating a part-of-speech tagged word corpus

The following is an excerpt from the brown corpus: `The/at-tl expense/nn and/cc time/nn involved/vbn are/ber astronomical/jj ./.` For example, treebank tag mappings are in `nltk_data/taggers/universal_tagset/en-ptb.map`.

## Creating a chunked phrase corpus

. O ends the sentence.

## Creating a categorized text corpus

The following are two excerpts from the movie\_reviews corpus: `movie_pos.txt`: `f the thin red line is flawed but it provokes .` Instead of `cat_pattern`, you could pass in a `cat_map`, which is a dictionary mapping a fileid argument to a list of category labels, as shown in the following code: `>>> reader = CategorizedPlaintextCorpusReader('.', r'movie_.*\.txt', cat_map={'movie_pos.txt': ['pos'], 'movie_neg.txt': ['neg']}) >>> reader.categories() ['neg', 'pos']` 65 [www.it-ebooks.info](http://www.it-ebooks.info) .

## Creating a categorized chunk corpus reader

The following code is found in `catchunked.py`: `from nltk.corpus.reader import CategorizedCorpusReader, ChunkedCorpusReader`

```

class CategorizedChunkedCorpusReader(CategorizedCorpusReader, ChunkedCorpusReader):
    def __init__(self, *args, **kwargs):
        CategorizedCorpusReader.__init__(self, *args, **kwargs)
        ChunkedCorpusReader.__init__(self, *args, **kwargs)
    def _resolve(self, fileids, categories):
        if fileids is not None and categories is not None:
            raise ValueError('Specify fileids or categories, not both')
        if categories is not None:
            return self.fileids(categories)
        else:
            return fileids
    All of the following methods call the corresponding function in ChunkedCorpusReader with the value returned from _resolve().

from nltk.corpus.reader import CategorizedCorpusReader, ConllCorpusReader, ConllChunkCorpusReader
class CategorizedConllChunkCorpusReader(CategorizedCorpusReader, ConllChunkCorpusReader):
    def __init__(self, *args, **kwargs):
        CategorizedCorpusReader.__init__(self, *args, **kwargs)
        ConllChunkCorpusReader.__init__(self, *args, **kwargs)
    def _resolve(self, fileids, categories):
        if fileids is not None and categories is not None:
            raise ValueError('Specify fileids or categories, not both')
        if categories is not None:
            return self.fileids(categories)
        else:
            return fileids
    All the following methods call the corresponding method of ConllCorpusReader with the value returned from _resolve().

```

## Lazy corpus loading

To speed up module import time when a corpus reader is defined, NLTK provides a `LazyCorpusLoader` class that can transform itself into your actual corpus reader as soon as you need it. The `LazyCorpusLoader` class requires two arguments: the name of the corpus and the corpus reader class, plus any other arguments needed to initialize the corpus reader class.

## Creating a custom corpus view

The following is the code found in corpus.py: from nltk.corpus.reader import PlaintextCorpusReader from nltk.corpus.reader.util import StreamBackedCorpusView class IgnoreHeadingCorpusView(StreamBackedCorpusView): def \_\_init\_\_(self, \*args, \*\*kwargs): StreamBackedCorpusView.\_\_init\_\_(self, \*args, \*\*kwargs) # open self.\_stream self.\_open() # skip the heading block self.read\_block(self.\_stream) # reset the start position to the current position in the stream self.\_filepos = [self.\_stream.tell()] class IgnoreHeadingCorpusReader(PlaintextCorpusReader): CorpusView = IgnoreHeadingCorpusView To demonstrate that this works as expected, here is code showing that the default PlaintextCorpusReader class finds four paragraphs, while our IgnoreHeadingCorpusReader class only has three paragraphs: >>> from nltk.corpus.reader import PlaintextCorpusReader >>> plain = PlaintextCorpusReader('.', ['heading\_text.txt']) >>> len(plain.paras()) 4 >>> from corpus import IgnoreHeadingCorpusReader >>> reader = IgnoreHeadingCorpusReader('.', ['heading\_text.txt']) >>> len(reader.paras()) 3 76 www.it-ebooks.info Chapter 3 How it works... In the Creating a part-of-speech tagged word corpus recipe, we discussed the default para\_block\_reader function of the TaggedCorpusReader class, which reads lines from a file until it finds a blank line, then returns those lines as a single paragraph token.

## Creating a MongoDB-backed corpus reader

The following is the code, which is found in mongoreader.py: import pymongo from nltk.data import LazyLoader from nltk.tokenize import TreebankWordTokenizer from nltk.util import AbstractLazySequence, LazyMap, LazyConcatenation class MongoDBLazySequence(AbstractLazySequence): def \_\_init\_\_(self, host='localhost', port=27017, db='test', collection='corpus', field='text'): self.conn = pymongo.MongoClient(host, port) self.collection = self.conn[db][collection] self.field = field def \_\_len\_\_(self): return self.collection.count() def iterate\_from(self, start): f = lambda d: d.get(self.field, "") return iter(LazyMap(f, self.collection.find(fields=[self.field], skip=start))) class MongoDBCorpusReader(object): def \_\_init\_\_(self, word\_tokenizer=TreebankWordTokenizer(), sent\_tokenizer=LazyLoader('tokenizers/punkt/PY3/english.pickle'), \*\*kwargs): self.\_seq = MongoDBLazySequence(\*\*kwargs) self.\_word\_tokenize = word\_tokenizer.tokenize self.\_sent\_tokenize = sent\_tokenizer.tokenize def text(self): return self.\_seq def words(self): return LazyConcatenation(LazyMap(self.\_word\_tokenize, self.text())) def sents(self): return LazyConcatenation(LazyMap(self.\_sent\_tokenize, self.text())) How it works... For example, if you had a db named website, with a collection named comments, whose documents had a field called comment, you could create a MongoDBCorpusReader class as follows: >>> reader = MongoDBCorpusReader(db='website', collection='comments', field='comment') You can also pass in custom instances for word\_tokenizer and sent\_tokenizer, as long as the objects implement the nltk.tokenize.TokenizerI interface by providing a tokenize(text) method.



# Corpus editing with file locking

These functions can be found in `corpus.py`, as follows:

```
import lockfile, tempfile, shutil
def append_line(fname, line):
    with lockfile.FileLock(fname):
        fp = open(fname, 'a+')
        fp.write(line)
        fp.write('\n')
        fp.close()
def remove_line(fname, line):
    with lockfile.FileLock(fname):
        tmp = tempfile.TemporaryFile()
        fp = open(fname, 'rw+')
        # write all lines from orig file, except if matches given line
        for l in fp:
            if l.strip() != line:
                tmp.write(l)
        # reset file pointers so entire files are copied
        fp.seek(0)
        tmp.seek(0)
        # copy tmp into fp, then truncate to remove trailing line(s)
        shutil.copyfileobj(tmp, fp)
        fp.truncate()
        fp.close()
        tmp.close()

```

The lock acquiring and releasing happens transparently when you do with `lockfile`. Once the lock is acquired, any other process that tries to access the file will have to wait until the lock is released.

## Introduction

[www.it-ebooks.info](http://www.it-ebooks.info) . The tag is a part-of-speech tag, and signifies whether the word is a noun, adjective, verb, and so on.

## Default tagging

Every subclass of `SequentialBackoffTagger` must implement the `choose_tag()` method, which takes three arguments: The list of tokens `f` The index of the current token whose tag we want to choose `f` The history, which is a list of the previous tags `f` `SequentialBackoffTagger` implements the `tag()` method, which calls the `choose_tag()` method of the subclass for each index in the tokens list while accumulating a history of the previously tagged tokens. `DefaultTagger` is most useful when you choose the most common part-of-speech tag.

## Training a unigram part-of-speech tagger

The tag with the highest frequency for a context is stored in the model. In other words, `UnigramTagger` is a context-based tagger whose context is a single word, or unigram.

## Combining taggers with backoff tagging

It allows you to chain taggers together so that if one tagger doesn't know how to tag a word, it can pass the word on to the next backoff tagger. But if it returns None, then the next tagger is tried, and so on until a tag is found, or else None is returned.

## Training and combining ngram taggers

Here's the code from tag\_util.py: `def backoff_tagger(train_sents, tagger_classes, backoff=None):` for `cls in tagger_classes:` `backoff = cls(train_sents, backoff=backoff)` `return backoff` And to use it, we can do the following: `>>> from tag_util import backoff_tagger >>> backoff = DefaultTagger('NN') >>> tagger = backoff_tagger(train_sents, [UnigramTagger, BigramTagger, TrigramTagger], backoff=backoff) >>> tagger.evaluate(test_sents)` 0.8806820634578028 So, we've gained almost 1% accuracy by including the BigramTagger and TrigramTagger subclasses in the backoff chain. `from nltk.tag import NgramTagger class QuadgramTagger(NgramTagger):` `def __init__(self, *args, **kwargs):` `NgramTagger.__init__(self, 4, *args, **kwargs)` This is essentially how BigramTagger and TrigramTagger are implemented: simple subclasses of NgramTagger that pass in the number of ngrams to look at in the history argument of the context() method.

## Creating a model of likely word tags

Then, we get the top 200 words from the FreqDist class by calling `fd.most_common()`, which obviously returns a list of the most common words and counts. `from nltk.probability import FreqDist, ConditionalFreqDist` `def word_tag_model(words, tagged_words, limit=200):` `fd = FreqDist(words)` `cfd = ConditionalFreqDist(tagged_words)` `most_freq = (word for word, count in fd.most_common(limit))` `return dict((word, cfd[word].max()) for word in most_freq)` And to use it with a UnigramTagger class, we can do the following: `>>> from tag_util import word_tag_model >>> from nltk.corpus import treebank >>> model = word_tag_model(treebank.words(), treebank.tagged_words()) >>> tagger = UnigramTagger(model=model) >>> tagger.evaluate(test_sents)` 0.559680552557738 An accuracy of almost 56% is ok, but nowhere near as good as the trained UnigramTagger.

## Tagging with regular expressions

The patterns shown in the following code can be found in tag\_util.py: `patterns = [ (r'^\d+$', 'CD'), (r'.*ing$', 'VBG'), # gerunds, i.e. wonderment (r'.*ful$', 'JJ') # i.e.`

## Affix tagging

Here's an example of four AffixTagger classes learning on 2 and 3 character prefixes and suffixes:

```
>>> pre3_tagger = AffixTagger(train_sents, affix_length=3) >>> pre3_tagger.evaluate(test_sents)
0.23587308439456076 >>> pre2_tagger = AffixTagger(train_sents, affix_length=2,
backoff=pre3_tagger) >>> pre2_tagger.evaluate(test_sents) 0.29786315562270665 >>> suf2_tagger
= AffixTagger(train_sents, affix_length=-2, backoff=pre2_tagger) >>>
suf2_tagger.evaluate(test_sents) 0.32467083962875026 >>> suf3_tagger = AffixTagger(train_sents,
affix_length=-3, backoff=suf2_tagger) >>> suf3_tagger.evaluate(test_sents) 0.3590761925318368
```

As you can see, the accuracy goes up each time. If a word is less than five characters, then None is returned as the tag.

## Training a Brill tagger

```
from nltk.tag import brill, brill_trainer
def train_brill_tagger(initial_tagger, train_sents, **kwargs):
    templates = [
        brill.Template(brill.Pos([-1])),    brill.Template(brill.Pos([1])),
        brill.Template(brill.Pos([-2])),    brill.Template(brill.Pos([2])),    brill.Template(brill.Pos([-2, -1])),
        brill.Template(brill.Pos([1, 2])),    brill.Template(brill.Pos([-3, -2, -1])),    brill.Template(brill.Pos([1, 2,
3])),    brill.Template(brill.Pos([-1]), brill.Pos([1])),    brill.Template(brill.Word([-1])),
        brill.Template(brill.Word([1])),    102 www.it-ebooks.info Chapter 4    brill.Template(brill.Word([-2])),
        brill.Template(brill.Word([2])),    brill.Template(brill.Word([-2, -1])),    brill.Template(brill.Word([1, 2])),
        brill.Template(brill.Word([-3, -2, -1])),    brill.Template(brill.Word([1, 2, 3])),
        brill.Template(brill.Word([-1]), brill.Word([1])),    ]
    trainer = brill_trainer.BrillTaggerTrainer(initial_tagger,
        templates, deterministic=True)
    return trainer.train(train_sents, **kwargs)
```

To use it, we can create our initial\_tagger from a backoff chain of NgramTagger classes, then pass that into the train\_brill\_tagger() function to get a BrillTagger back. Found 9869 useful rules.

## Training the TnT tagger

However, don't assume that accuracy will not change if you decrease N; experiment with your own data to be sure. You must explicitly call the train() method after you've created it.

## Using WordNet for tagging

The WordNetTagger class defined in the following code can be found in taggers.py:

```
from nltk.tag
import SequentialBackoffTagger from nltk.corpus import wordnet from nltk.probability import FreqDist
class WordNetTagger(SequentialBackoffTagger):
    """
    >>> wt = WordNetTagger()
    >>> wt.tag(['food', 'is', 'great'])
    [('food', 'NN'), ('is', 'VB'), ('great', 'JJ')]
    """
    def __init__(self, *args, **kwargs):
        SequentialBackoffTagger.__init__(self, *args, **kwargs)
        self.wordnet_tag_map = {
            'n': 'NN',
            's': 'JJ',
            'a': 'JJ',
            'r': 'RB',
            'v': 'VB'
        }
    def choose_tag(self, tokens, index, history):
        word = tokens[index]
        fd = FreqDist()
        for synset in wordnet.synsets(word):
            fd[synset.pos()] += 1
        return self.wordnet_tag_map.get(fd.max())

108 www.it-ebooks.info
```

Chapter 4 Another way the FreqDist API has changed between NLTK2 and NLTK3 is that the inc() method has been removed. It's a very restricted set of possible tags, and many words have multiple Synsets with different part-of-speech tags, but this information can be useful for tagging unknown words.

## Tagging proper names

The following code can be found in taggers.py:

```
from nltk.tag import SequentialBackoffTagger from nltk.corpus import names
class NamesTagger(SequentialBackoffTagger):
    def __init__(self, *args, **kwargs):
        SequentialBackoffTagger.__init__(self, *args, **kwargs)
        self.name_set = set([n.lower() for n in names.words()])
    def choose_tag(self, tokens, index, history):
        word = tokens[index]
        if word.lower() in self.name_set:
            return 'NNP'
        else:
            return None
```

How it works... If it isn't, we return None, so the next tagger in the chain can tag the word.

## Classifier-based tagging

The ClassifierBasedPOSTagger class is a subclass of ClassifierBasedTagger that implements a feature detector that combines many of the techniques of the previous taggers into a single feature set. For example, to use a MaxentClassifier, you'd do the following:

```
>>> from nltk.classify import MaxentClassifier
>>> me_tagger = ClassifierBasedPOSTagger(train=train_sents,
classifier_builder=MaxentClassifier.train)
>>> me_tagger.evaluate(test_sents)
0.9258363911072739
```

The MaxentClassifier class takes even longer to train than NaiveBayesClassifier.

## Training a tagger with NLTK-Trainer

Found 1304 useful rules. As we did earlier, we can boost the accuracy a bit by using a default tagger:

```
$ python train_tagger.py treebank --no-pickle --default NN --fraction 0.75 --sequential u loading
treebank 3914 tagged sents, training on 2936 training <class 'nltk.tag.sequential.UnigramTagger'>
tagger with backoff <DefaultTagger: tag=NN> evaluating UnigramTagger accuracy: 0.873462 116
www.it-ebooks.info Chapter 4 Now, let's try adding a bigram tagger and trigram tagger: $ python
train_tagger.py treebank --no-pickle --default NN --fraction 0.75 --sequential ubt loading treebank
3914 tagged sents, training on 2936 training <class 'nltk.tag.sequential.UnigramTagger'> tagger with
backoff <DefaultTagger: tag=NN> training <class 'nltk.tag.sequential.BigramTagger'> tagger with
backoff <UnigramTagger: size=8709> training <class 'nltk.tag.sequential.TrigramTagger'> tagger with
backoff <BigramTagger: size=1836> evaluating TrigramTagger accuracy: 0.879012 The
PYTHONHASHSEED environment variable has been omitted for clarity.
```

## Introduction

The idea is that meaningful phrases can be extracted from a sentence by looking for particular patterns of part-of-speech tags. [www.it-ebooks.info](http://www.it-ebooks.info) .

## Chunking and chinking with regular expressions

```
}<VB.*>{ ... {<DT><NN.*><.*>*<NN.*>} ...
```

## Merging and splitting chunks with regular expressions

A `SplitRule` class will split a chunk into two chunks based on the specified split pattern. A `MergeRule` class is specified by flipping the curly braces, and will join chunks where the end of the first chunk matches the left pattern and the beginning of the next chunk matches the right pattern.

## Expanding and removing chunks with regular expressions

These rules are as follows: `ExpandLeftRule`: Add unchunked (chink) words to the left of a chunk `f` `ExpandRightRule`: Add unchunked (chink) words to the right of a chunk `f` `UnChunkRule`: Unchunk any matching chunk `f` 133 [www.it-ebooks.info](http://www.it-ebooks.info) Extracting Chunks How to do it... That's because the final `UnChunkRule` undid the chunk created by the previous rules.

## Partial parsing with regular expressions

{<IN>} # chunk preposition ... Of the chunks the chunker argument was able to guess, precision tells you how many were correct and recall tells you how well the chunker did at finding correct chunks compared to how many total chunks there were.

## Training a tagger-based chunker

Here's the code from chunkers.py: from nltk.chunk import ChunkParserI from nltk.chunk.util import tree2conlltags, conlltags2tree from nltk.tag import UnigramTagger, BigramTagger from tag\_util import backoff\_tagger 139 www.it-ebooks.info Extracting Chunks def conll\_tag\_chunks(chunk\_sents): tagged\_sents = [tree2conlltags(tree) for tree in chunk\_sents] return [[(t, c) for (w, t, c) in sent] for sent in tagged\_sents] class TagChunker(ChunkParserI): def \_\_init\_\_(self, train\_chunks, tagger\_classes=[UnigramTagger, BigramTagger]): train\_sents = conll\_tag\_chunks(train\_chunks) self.tagger = backoff\_tagger(train\_sents, tagger\_classes) def parse(self, tagged\_sent): if not tagged\_sent: return None (words, tags) = zip(\*tagged\_sent) chunks = self.tagger.tag(tags) wtc = zip(words, chunks) return conlltags2tree([(w,t,c) for (w,(t,c)) in wtc]) Once we have our trained TagChunker, we can then evaluate the ChunkScore class the same way we did for the RegexpParser class in the previous recipes: >>> from chunkers import TagChunker >>> from nltk.corpus import treebank\_chunk >>> train\_chunks = treebank\_chunk.chunked\_sents()[:3000] >>> test\_chunks = treebank\_chunk.chunked\_sents()[3000:] >>> chunker = TagChunker(train\_chunks) >>> score = chunker.evaluate(test\_chunks) >>> score.accuracy() 0.9732039335251428 >>> score.precision() 0.9166534370535006 >>> score.recall() 0.9465573770491803 Pretty darn accurate! Training a chunker is clearly a great alternative to manually specified grammars and regular expressions. The tags are then tagged by the tagger to get IOB tags, which are then recombined with the words and part-of-speech tags to create 3-tuples we can pass to conlltags2tree() to return a final Tree.

## Classification-based chunking

As a subclass of `ChunkerParserI`, it implements the `parse()` method, which converts the `((w, t), c)` tuples produced by the internal tagger into `Trees` using `conlltags2tree()`:

```
class ClassifierChunker(ChunkParserI):
    def __init__(self, train_sents,
                  feature_detector=prev_next_pos_iob, **kwargs):
        if not feature_detector:
            feature_detector = self.feature_detector
        train_chunks = chunk_trees2train_chunks(train_sents)
        self.tagger = ClassifierBasedTagger(train=train_chunks,
                                           feature_detector=feature_detector, **kwargs)

    def parse(self, tagged_sent):
        if not tagged_sent:
            return None
        chunks = self.tagger.tag(tagged_sent)
        return conlltags2tree([(w,t,c) for ((w,t),c) in chunks])
```

Using the same `train_chunks` and `test_chunks` from the `treebank_chunk` corpus in the previous recipe, we can evaluate this code from `chunkers.py`:

```
>>> from chunkers import ClassifierChunker
>>> chunker = ClassifierChunker(train_chunks)
>>> score = chunker.evaluate(test_chunks)
>>> score.accuracy()
0.9721733155838022
>>> score.precision()
0.9258838793383068
>>> score.recall()
0.9359016393442623
```

Compared to the `TagChunker` class, all the scores have gone up a bit. To give the classifier as much information as we can, this feature set contains the current, previous, and next word and part-of-speech tag, along with the previous IOB tag:

```
def prev_next_pos_iob(tokens, index, history):
    word, pos = tokens[index]
    if index == 0:
        prevword, prevpos, previob = ('<START>')*3
    else:
        prevword, prevpos = tokens[index-1]
        previob = history[index-1]
```

143 [www.it-ebooks.info](http://www.it-ebooks.info)

```
def nextword, nextpos = ('<END>')*2
    else:
        nextword, nextpos = tokens[index+1]
    feats = {
        'word': word,
        'pos': pos,
        'nextword': nextword,
        'nextpos': nextpos,
        'prevword': prevword,
        'prevpos': prevpos,
        'previob': previob
    }
    return feats
```

Now, we can define the `ClassifierChunker` class, which uses an internal `ClassifierBasedTagger` with features extracted using `prev_next_pos_iob()` and training sentences from `chunk_trees2train_chunks()`.

## Extracting named entities

If we get the `sub_leaves()`, we can see that `Pierre Vinken` is correctly combined into a single named entity:

```
>>> subleaves(ne_chunk(treebank_chunk.tagged_sents()[0], binary=True), 'NE')
[('Pierre', 'NNP'), ('Vinken', 'NNP')]
```

148 [www.it-ebooks.info](http://www.it-ebooks.info)

Now, all named entities will be tagged with `NE`:

```
>>> ne_chunk(treebank_chunk.tagged_sents()[0], binary=True)
Tree('S', [Tree('NE', [('Pierre', 'NNP'), ('Vinken', 'NNP')]), (',', ','), ('61', 'CD'), ('years', 'NNS'), ('old', 'JJ'), ('.', '.'), ('will', 'MD'), ('join', 'VB'), ('the', 'DT'), ('board', 'NN'), ('as', 'IN'), ('a', 'DT'), ('nonexecutive', 'JJ'), ('director', 'NN'), ('Nov.', 'NNP'), ('29', 'CD'), ('.', '.')])
```

So, `binary` in this case means that an arbitrary chunk either is or is not a named entity.

## Extracting proper noun chunks

This class can be found in chunkers.py: from nltk.chunk import ChunkParserI from nltk.chunk.util import conlltags2tree from nltk.corpus import names 149 www.it-ebooks.info Extracting Chunks class PersonChunker(ChunkParserI): def \_\_init\_\_(self): self.name\_set = set(names.words()) def parse(self, tagged\_sent): iobs = [] in\_person = False for word, tag in tagged\_sent: if word in self.name\_set and in\_person: iobs.append((word, tag, 'I-PERSON')) elif word in self.name\_set: iobs.append((word, tag, 'B-PERSON')) in\_person = True else: iobs.append((word, tag, 'O')) in\_person = False return conlltags2tree(iobs) The PersonChunker class iterates over the tagged sentence, checking whether each word is in its name\_set (constructed from the names corpus). is the name of a month.

## Extracting location chunks

The helper method iob\_locations() is where the IOB LOCATION tags are produced, and the parse() method converts these IOB tags into a Tree: from nltk.chunk import ChunkParserI from nltk.chunk.util import conlltags2tree from nltk.corpus import gazetteers class LocationChunker(ChunkParserI): def \_\_init\_\_(self): self.locations = set(gazetteers.words()) self.lookahead = 0 for loc in self.locations: nwords = loc.count(' ') if nwords > self.lookahead: self.lookahead = nwords 151 www.it-ebooks.info Extracting Chunks def iob\_locations(self, tagged\_sent): i = 0 l = len(tagged\_sent) inside = False while i < l: word, tag = tagged\_sent[i] j = i + 1 k = j + self.lookahead nextwords, nexttags = [], [] loc = False while j < k: if ''.join([word] + nextwords) in self.locations: if inside: yield word, tag, 'I-LOCATION' else: yield word, tag, 'B-LOCATION' for nword, ntag in zip(nextwords, nexttags): yield nword, ntag, 'I-LOCATION' loc, inside = True, True i = j break if j < l: nextword, nexttag = tagged\_sent[j] nextwords.append(nextword) nexttags.append(nexttag) j += 1 else: break if not loc: inside = False i += 1 yield word, tag, 'O' def parse(self, tagged\_sent): iobs = self.iob\_locations(tagged\_sent) return conlltags2tree(iobs) 152 www.it-ebooks.info Chapter 5 We can use the LocationChunker class to parse the following sentence into two locations?San Francisco CA is cold compared to San Jose CA: >>> from chunkers import LocationChunker >>> t = loc.parse([('San', 'NNP'), ('Francisco', 'NNP'), ('CA', 'NNP'), ('is', 'BE'), ('cold', 'JJ'), ('compared', 'VBD'), ('to', 'TO'), ('San', 'NNP'), ('Jose', 'NNP'), ('CA', 'NNP')]) >>> sub\_leaves(t, 'LOCATION') [[('San', 'NNP'), ('Francisco', 'NNP'), ('CA', 'NNP')], [('San', 'NNP'), ('Jose', 'NNP'), ('CA', 'NNP')]] And the result is that we get two LOCATION chunks, just as expected. The gazetteers corpus is a WordListCorpusReader class that contains the following location words: Country names f U.S.

## Training a named entity chunker



Using the `ieertree2conlltags()` and `ieer_chunked_sents()` functions in `chunkers.py`, we can create named entity chunk trees from the `ieer` corpus to train the `ClassifierChunker` class created in the Classification-based chunking recipe: `import nltk.tag from nltk.chunk.util import conlltags2tree from nltk.corpus import ieer def ieertree2conlltags(tree, tag=nltk.tag.pos_tag): words, ents = zip(*tree.pos()) iobs = [] prev = None for ent in ents: if ent == tree.label(): iobs.append('O') prev = None elif prev == ent: iobs.append('I-%s' % ent) else: iobs.append('B-%s' % ent) prev = ent words, tags = zip(*tag(words)) return zip(words, tags, iobs) def ieer_chunked_sents(tag=nltk.tag.pos_tag): for doc in ieer.parsed_docs(): tagged = ieertree2conlltags(doc.text, tag) yield conlltags2tree(tagged)` 154 [www.it-ebooks.info](http://www.it-ebooks.info) Chapter 5 We'll use 80 out of 94 sentences for training, and the rest for testing. Once we have all the IOB tags, then we can get the part-of-speech tags of all the words and join the words, part-of-speech tags, and IOB tags into 3-tuples using `zip()`.

## Training a chunker with NLTK-Trainer

Training classifier (47000 instances) training NaiveBayes classifier evaluating ClassifierChunker  
 ChunkParse score: IOB Accuracy: 88.3% Precision: 40.9% Recall: 50.5%  
 F-Measure: 45.2% Training on a custom corpus If you have a custom corpus that you want to use for training a chunker, you can do that by passing in the path to the corpus and the classname of a corpus reader in the `--reader` argument. Here's an example of running `train_chunker.py` on `treebank_chunk`:  
 \$ python train\_chunker.py treebank\_chunk loading treebank\_chunk 4009 chunks, training on 4009 training ub TagChunker evaluating TagChunker ChunkParse score: IOB Accuracy: 97.0% Precision: 90.8% Recall: 93.9% F-Measure: 92.3% dumping TagChunker to /Users/jacob/nltk\_data/chunkers/treebank\_chunk\_ub.pickle 156 [www.it-ebooks.info](http://www.it-ebooks.info) Chapter 5 Just like with `train_tagger.py`, we can use the `--no-pickle` argument to skip saving a pickled chunker, and the `--fraction` argument to limit the training set and evaluate the chunker against a test set:  
 \$ python train\_chunker.py treebank\_chunk --no-pickle --fraction 0.75 loading treebank\_chunk 4009 chunks, training on 3007 training ub TagChunker evaluating TagChunker ChunkParse score: IOB Accuracy: 97.3% Precision: 91.6% Recall: 94.6% F-Measure: 93.1% The score output you see is what you get when you print a `ChunkScore` object.

## Introduction

[www.it-ebooks.info](http://www.it-ebooks.info) . The tree transforms give you ways to modify and flatten deep parse trees.

## Filtering insignificant words from a sentence

It defaults to filtering out any tags that end with DT or CC: `def filter_insignificant(chunk, tag_suffixes=['DT', 'CC']):` `good = []` 164 [www.it-ebooks.info](http://www.it-ebooks.info) Chapter 6 `for word, tag in chunk:` `ok = True` `for suffix in tag_suffixes:` `if tag.endswith(suffix):` `ok = False` `break` `if ok:` `good.append((word, tag))` `return good` And now we can use it on the part-of-speech tagged version of the terrible movie: `>>> from transforms import filter_insignificant >>> filter_insignificant([('the', 'DT'), ('terrible', 'JJ'), ('movie', 'NN')])` `[('terrible', 'JJ'), ('movie', 'NN')]` As you can see, the word the is eliminated from the chunk. Looking through the treebank corpus for stopwords yields the following table of insignificant words and tags: Word Tag a DT all PDT an DT and CC or CC that WDT the DT Other than CC, all the tags end with DT.

## Correcting verb forms

`def correct_verbs(chunk):` `vbidx = first_chunk_index(chunk, tagstartswith('VB'))` `# if no verb found,` `do nothing` `if vbidx is None:` `return chunk` `verb, vntag = chunk[vbidx]` `nnpred = tagstartswith('NN')` `# find nearest noun to the right of verb` `nnidx = first_chunk_index(chunk, nnpred, start=vbidx+1)` `# if no noun found to right, look to the left` `if nnidx is None:` `nnidx = first_chunk_index(chunk, nnpred, start=vbidx-1, step=-1)` `# if no noun found, do nothing` `if nnidx is None:` `return chunk` 167 [www.it-ebooks.info](http://www.it-ebooks.info) Transforming Chunks and Trees `noun, nntag = chunk[nnidx]` `# get correct verb form and insert into chunk` `if nntag.endswith('S'):` `chunk[vbidx] = plural_verb_forms.get((verb, vntag), (verb, vntag))` `else:` `chunk[vbidx] = singular_verb_forms.get((verb, vntag), (verb, vntag))` `return chunk` When we call the preceding function on a part-of-speech tagged is our children learning chunk, we get back the correct form, are our children learning. `def tagstartswith(prefix):` `def f(wt):` `return wt[1].startswith(prefix)` `return f` The `tagstartswith()` function takes a tag prefix, such as NN, and returns a predicate function that will take a (word, tag) tuple and return True if the tag starts with the given prefix.

## Swapping verb phrases

It uses the `first_chunk_index()` function defined in the previous recipe to find the verb to pivot around. `def swap_verb_phrase(chunk):` `def vbpred(wt):` `word, tag = wt` `return tag != 'VBG' and tag.startswith('VB') and len(tag) > 2` `vbidx = first_chunk_index(chunk, vbpred)` `if vbidx is None:` `return chunk` `return chunk[vbidx+1:] + chunk[:vbidx]` Now we can see how it works on the part-of-speech tagged phrase the book was great: `>>> swap_verb_phrase([('the', 'DT'), ('book', 'NN'), ('was', 'VBD'), ('great', 'JJ')])` `[('great', 'JJ'), ('the', 'DT'), ('book', 'NN')]` And the result is great the book.

## Swapping noun cardinals

It swaps any cardinal that occurs immediately after a noun with the noun so that the cardinal occurs immediately before the noun. It uses a helper function, `tag_equals()`, which is similar to `tag_startswith()`, but in this case, the function it returns does an equality comparison with the given tag: `def tag_equals(tag): def f(wt): return wt[1] == tag return f` Now we can define `swap_noun_cardinal()`: `def swap_noun_cardinal(chunk): cidx = first_chunk_index(chunk, tag_equals('CD')) # cidx must be > 0 and there must be a noun immediately before it if not cidx or not chunk[cidx-1][1].startswith('NN'): return chunk noun, nntag = chunk[cidx-1] chunk[cidx-1] = chunk[cidx] chunk[cidx] = noun, nntag return chunk` Let's try it on a date, such as Dec 10, and another common phrase, the top 10.

## Swapping infinitive phrases

The `swap_infinitive_phrase()` function, defined in `transforms.py`, will return a chunk that swaps the portion of the phrase after the IN word with the portion before the IN word: `def swap_infinitive_phrase(chunk): def inpred(wt): word, tag = wt return tag == 'IN' and word != 'like' inidx = first_chunk_index(chunk, inpred) if inidx is None: return chunk nnidx = first_chunk_index(chunk, tag_startswith('NN'), start=inidx, step=-1) or 0 return chunk[:nnidx] + chunk[inidx+1:] + chunk[nnidx:inidx]` The function can now be used to transform `book of recipes` into `recipes book`: `>>> from transforms import swap_infinitive_phrase >>> swap_infinitive_phrase([('book', 'NN'), ('of', 'IN'), ('recipes', 'NNS')])` `[('recipes', 'NNS'), ('book', 'NN')]` How it works... Instead, we want to insert `recipes` before the noun `book` but after the adjective `delicious`, hence the need to find the `nnidx` occurring before the `inidx`.

## Singularizing plural nouns

The `transforms.py` script defines a function called `singularize_plural_noun()` which will depluralize a plural noun (tagged with `NNS`) that is followed by another noun: `def singularize_plural_noun(chunk): nnsidx = first_chunk_index(chunk, tag_equals('NNS')) if nnsidx is not None and nnsidx+1 < len(chunk) and chunk[nnsidx+1][1][2] == 'NN': noun, nntag = chunk[nnsidx] chunk[nnsidx] = (noun.rstrip('s'), nntag.rstrip('S')) return chunk` And using it on `recipes book`, we get the more correct form, `recipe book`. We can do another transform to correct these improper plural nouns.

## Chaining chunk transformations

It calls each transform function on the chunk, one at a time, and returns the final chunk: `def transform_chunk(chunk, chain=[filter_insignificant, swap_verb_phrase, swap_infinitive_phrase, singularize_plural_noun], trace=0):` for f in chain: chunk = f(chunk) if trace: print f.\_\_name\_\_, ':', chunk return chunk Using it on the phrase the book of recipes is delicious, we get delicious recipe book: `>>> from transforms import transform_chunk >>> transform_chunk([(('the', 'DT'), ('book', 'NN'), ('of', 'IN'), ('recipes', 'NNS'), ('is', 'VBZ'), ('delicious', 'JJ'))], [('delicious', 'JJ'), ('recipe', 'NN'), ('book', 'NN')])` 174 [www.it-ebooks.info](http://www.it-ebooks.info) Chapter 6 How it works... The `transform_chunk()` function defaults to chaining the following functions in the given order: `filter_insignificant()` f `swap_verb_phrase()` f `swap_infinitive_phrase()` f `singularize_plural_noun()` f Each function transforms the chunk that results from the previous function, starting with the original chunk.

## Converting a chunk tree to text

This is implemented in the `chunk_tree_to_sent()` function found in `transforms.py`: `import re punct_re = re.compile(r'\s([\.,\;\?])')` `def chunk_tree_to_sent(tree, concat=' '):` s = concat.join([w for w, t in tree.leaves()]) return re.sub(punct\_re, r'\g<1>', s) Using `chunk_tree_to_sent()` results in a cleaner sentence, with no space before each punctuation mark: `>>> from transforms import chunk_tree_to_sent >>> chunk_tree_to_sent(tree)` 'Pierre Vinken, 61 years old, will join the board as a nonexecutive director Nov. To correct the extra spaces in front of the punctuation, we create a regular expression, `punct_re`, that will match a space followed by any of the known punctuation characters.

## Flattening a deep tree

It uses a helper function, `flatten_childtrees()`, to do most of the work: `from nltk.tree import Tree def flatten_childtrees(trees):` children = [] for t in trees: if t.height() < 3: children.extend(t.pos()) elif t.height() == 3: children.append(Tree(t.label(), t.pos())) else: children.extend(flatten\_childtrees([c for c in t])) return children `def flatten_deeptree(tree):` return Tree(tree.label(), flatten\_childtrees([c for c in tree])) We can use it on the first parsed sentence of the treebank corpus to get a flatter tree: `>>> from nltk.corpus import treebank >>> from transforms import flatten_deeptree >>> flatten_deeptree(treebank.parsed_sents()[0])` Tree('S', [Tree('NP', [('Pierre', 'NNP'), ('Vinken', 'NNP')]), (',', ','), Tree('NP', [('61', 'CD'), ('years', 'NNS')]), ('old', 'JJ'), (',', ','), ('will', 'MD'), ('join', 'VB'), Tree('NP', [('the', 'DT'), ('board', 'NN')]), ('as', 'IN'), Tree('NP', [('a', 'DT'), ('nonexecutive', 'JJ'), ('director', 'NN')]), Tree('NP-TMP', [('Nov.', 'NNP'), ('29', 'CD')]), ('.', '.')]) 178 [www.it-ebooks.info](http://www.it-ebooks.info) Chapter 6 The result is a much flatter Tree that only includes NP phrases. The main difference is that the rightmost NP Tree is separated into two subtrees above, one of them named NP-TMP.

## Creating a shallow tree

29 a nonexecutive director The shallow\_tree() function defined in transforms.py eliminates all the nested subtrees, keeping only the top subtree labels: from nltk.tree import Tree def shallow\_tree(tree): children = [] for t in tree: if t.height() < 3: children.extend(t.pos()) else: children.append(Tree(t.label(), t.pos())) return Tree(tree.label(), children) Using it on the first parsed sentence in treebank results in a Tree with only two subtrees: >>> from transforms import shallow\_tree >>> shallow\_tree(treebank.parsed\_sents()[0]) Tree('S', [Tree('NP-SBJ', [('Pierre', 'NNP'), ('Vinken', 'NNP'), (',', ','), ('61', 'CD'), ('years', 'NNS'), ('old', 'JJ'), (',', ',')]), Tree('VP', [(('will', 'MD'), ('join', 'VB'), ('the', 'DT'), ('board', 'NN'), ('as', 'IN'), ('a', 'DT'), ('nonexecutive', 'JJ'), ('director', 'NN'), ('Nov.', 'NNP'), ('29', 'CD'))], (',', ','))]) 182 www.it-ebooks.info . In this recipe, we'll keep only the highest level subtrees instead.

## Converting tree labels

It returns a new Tree with all matching labels replaced based on the values in the mapping: from nltk.tree import Tree def convert\_tree\_labels(tree, mapping): children = [] for t in tree: if isinstance(t, Tree): children.append(convert\_tree\_labels(t, mapping)) else: children.append(t) label = mapping.get(tree.label(), tree.label()) return Tree(label, children) 184 www.it-ebooks.info Chapter 6 Using the mapping table we saw earlier, we can pass it in as a dict to convert\_tree\_labels() and convert the first parsed sentence from treebank: >>> from transforms import convert\_tree\_labels >>> mapping = {'NP-SBJ': 'NP', 'NP-TMP': 'NP'} >>> convert\_tree\_labels(treebank.parsed\_sents()[0], mapping) Tree('S', [Tree('NP', [Tree('NP', [Tree('NP', ['Pierre']), Tree('NP', ['Vinken'])]), Tree(',', [',']), Tree('ADJP', [Tree('NP', [Tree('CD', ['61']), Tree('NNS', ['years'])]), Tree('JJ', ['old'])]), Tree(',', [','])]), Tree('VP', [Tree('MD', ['will']), Tree('VP', [Tree('VB', ['join']), Tree('NP', [Tree('DT', ['the']), Tree('NN', ['board'])]), Tree('PP-CLR', [Tree('IN', ['as']), Tree('NP', [Tree('DT', ['a']), Tree('JJ', ['nonexecutive']), Tree('NN', ['director'])])]), Tree('NP', [Tree('NP', ['Nov.']), Tree('CD', ['29'])])])]), Tree('.', ['.'])])])]) As you can see in the following diagram, the NP-\* subtrees have been replaced with NP subtrees: S NP VP . Let's convert both of those to NP.

## Introduction

The text can either be one label or another, but not both, whereas a multi-label classifier can assign one or more labels to a piece of text. www.it-ebooks.info .

## Bag of words feature extraction

The `bag_of_bigrams_words()` function found in `featx.py` will return a dict of all words along with the 200 most significant bigrams: `from nltk.collocations import BigramCollocationFinder from nltk.metrics import BigramAssocMeasures def bag_of_bigrams_words(words, score_fn=BigramAssocMeasures.chi_sq, n=200): bigram_finder = BigramCollocationFinder.from_words(words) bigrams = bigram_finder.nbest(score_fn, n) return bag_of_words(words + bigrams)` The bigrams will be present in the returned dict as `(word1, word2)` and will have the value as `True`. The idea is to convert a list of words into a dict, where each word becomes a key with the value `True`.

## Training a Naive Bayes classifier

Getting a fair sample should eliminate this possible bias: `import collections def label_feats_from_corpus(corp, feature_detector=bag_of_words): label_feats = collections.defaultdict(list) for label in corp.categories(): for fileid in corp.fileids(categories=[label]): feats = feature_detector(corp.words(fileids=[fileid])) label_feats[label].append(feats) return label_feats` Once we can get a mapping of label | feature sets, we want to construct a list of labeled training instances and testing instances. In our case, the feature value will always be `True`: `>>> nb_classifier.most_informative_features(n=5) [('magnificent', True), ('outstanding', True), ('insulting', True), ('vulnerable', True), ('ludicrous', True)]` 195 [www.it-ebooks.info](http://www.it-ebooks.info) Text Classification The `show_most_informative_features()` method will print out the results from `most_informative_features()` and will also include the probability of a feature pair belonging to each label: `>>> nb_classifier.show_most_informative_features(n=5)` Most Informative Features `magnificent = True pos : neg = 15.0 : 1.0 outstanding = True pos : neg = 13.6 : 1.0 insulting = True neg : pos = 13.0 : 1.0 vulnerable = True pos : neg = 12.3 : 1.0 ludicrous = True neg : pos = 11.8 : 1.0` The informativeness, or information gain, of each feature pair is based on the prior probability of the feature pair occurring for each label.

## Training a decision tree classifier

We pass `binary=True` because all of our features are binary: either the word is present or it's not. The default value is 100, which means that classification may require up to 100 decisions before reaching a leaf node.

## Training a maximum entropy classifier

The only difference between these two algorithms that really matters is that `gis` is much faster than `iis`. In other words, if the word `worst` is found in the feature set, then there's a strong possibility that the text should be classified `neg`.

## Training scikit-learn classifiers

Here's the complete class code, minus all comments, docstrings, and most imports: from `sklearn.feature_extraction` import `DictVectorizer` from `sklearn.preprocessing` import `LabelEncoder`

```
class SklearnClassifier(ClassifierI):
    def __init__(self, estimator, dtype=float, sparse=True):
        self._clf = estimator
        self._encoder = LabelEncoder()
        self._vectorizer = DictVectorizer(dtype=dtype, sparse=sparse)

    def batch_classify(self, featuresets):
        X = self._vectorizer.transform(featuresets)
        classes = self._encoder.classes_
        return [classes[i] for i in self._clf.predict(X)]

    def batch_prob_classify(self, featuresets):
        X = self._vectorizer.transform(featuresets)
        y_proba_list = self._clf.predict_proba(X)
        return [self._make_probdist(y_proba) for y_proba in y_proba_list]

    def labels(self):
        return list(self._encoder.classes_)

    def train(self, labeled_featuresets):
        X, y = list(zip(*labeled_featuresets))
        self._vectorizer.fit_transform(X)
        y = self._encoder.fit_transform(y)
        self._clf.fit(X, y)
        return self

    def _make_probdist(self, y_proba):
        classes = self._encoder.classes_
        return DictionaryProbDist(dict((classes[i], p) for i, p in enumerate(y_proba)))
```

The class is initialized with an estimator, which is the algorithm we pass in, such as `MultinomialNB`.

```
>>> from sklearn.linear_model import LogisticRegression
>>> sk_classifier = SklearnClassifier(LogisticRegression())
<SklearnClassifier(LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
    intercept_scaling=1, penalty='l2', random_state=None, tol=0.0001))>
>>> sk_classifier.train(train_feats)
>>> accuracy(sk_classifier, test_feats)
0.892
```

Again, we see that the `sklearn` algorithm has better performance than `NLTK's MaxentClassifier`, which only had 72.2% accuracy.

## Measuring precision and recall of a classifier

The `precision_recall()` function in `classification.py` looks like this:

```
import collections
from nltk import metrics

def precision_recall(classifier, testfeats):
    refsets = collections.defaultdict(set)
    testsets = collections.defaultdict(set)
    for i, (feats, label) in enumerate(testfeats):
        refsets[label].add(i)
        observed = classifier.classify(feats)
        testsets[observed].add(i)
    precisions = {}
    recalls = {}
    for label in classifier.labels():
        precisions[label] = metrics.precision(refsets[label], testsets[label])
        recalls[label] = metrics.recall(refsets[label], testsets[label])
    return precisions, recalls
```

This function takes two arguments: The trained classifier and Labeled test features, also known as a gold standard. These are the same arguments you pass to `accuracy()`. Two of the most common are precision and recall.

## Calculating high information words

We can do this using the `high_information_words()` function in `featx.py`: `from nltk.metrics import BigramAssocMeasures from nltk.probability import FreqDist, ConditionalFreqDist def high_information_words(labelled_words, score_fn=BigramAssocMeasures.chi_sq, min_score=5): word_fd = FreqDist() label_word_fd = ConditionalFreqDist() for label, words in labelled_words: for word in words: word_fd[word] += 1 label_word_fd[label][word] += 1 n_xx = label_word_fd.N() high_info_words = set() for label in label_word_fd.conditions(): n_xi = label_word_fd[label].N() word_scores = collections.defaultdict(int) 214 www.it-ebooks.info Chapter 7 for word, n_ii in label_word_fd[label].items(): n_ix = word_fd[word] score = score_fn(n_ii, (n_ix, n_xi), n_xx) word_scores[word] = score bestwords = [word for word, score in word_scores.items() if score >= min_score] high_info_words |= set(bestwords) return high_info_words` It takes one argument from a list of two tuples of the form `[(label, words)]` where `label` is the classification label, and `words` is a list of words that occur under that label. Once we have the high information words, we use the feature detector function `bag_of_words_in_set()`, also found in `featx.py`, which will let us filter out all low information words.

## Combining classifiers with voting

In the `classification.py` module, there is a `MaxVoteClassifier` class: `import itertools from nltk.classify import ClassifierI from nltk.probability import FreqDist class MaxVoteClassifier(ClassifierI): def __init__(self, *classifiers): self._classifiers = classifiers self._labels = sorted(set(itertools.chain(*[c.labels() for c in classifiers]))) def labels(self): return self._labels def classify(self, feats): counts = FreqDist() for classifier in self._classifiers: counts[classifier.classify(feats)] += 1 return counts.max()` 219 www.it-ebooks.info Text Classification To create it, you pass in a list of classifiers that you want to combine. The `MaxVoteClassifier` class iterates over its classifiers and calls `classify()` on each of them, recording their label as a vote in a `FreqDist` variable.

## Classifying with multiple binary classifiers



The `train_binary_classifiers()` function in `classification.py` takes a training function, a list of multi-label feature sets, and a set of possible labels to return a dict of label : binary classifier:

```
def train_binary_classifiers(trainf, labelled_feats, labelset):
    pos_feats = collections.defaultdict(list)
    neg_feats = collections.defaultdict(list)
    classifiers = {}
    for feat, labels in labelled_feats:
        for label in labels:
            pos_feats[label].append(feat)
        for label in labelset - set(labels):
            neg_feats[label].append(feat)
    postrain = [(feat, label) for feat in pos_feats[label]]
    negtrain = [(feat, '!%' % label) for feat in neg_feats[label]]
    classifiers[label] = trainf(postrain + negtrain)
    return classifiers
```

To use this function, we need to provide a training function that takes a single argument, which is the training data.

```
from nltk.classify import MultiClassifierI
class MultiBinaryClassifier(MultiClassifierI):
    def __init__(self, *label_classifiers):
        self._label_classifiers = dict(label_classifiers)
        self._labels = sorted(self._label_classifiers.keys())
    def labels(self):
        return self._labels
    def classify(self, feats):
        lbls = set()
        for label, classifier in self._label_classifiers.items():
            if classifier.classify(feats) == label:
                lbls.add(label)
        return lbls
```

Now we can construct this class using the binary classifiers we just created:

```
>>> from classification import MultiBinaryClassifier
>>> multi_classifier = MultiBinaryClassifier(*classifiers.items())
```

To evaluate this classifier, we can use precision and recall, but not accuracy.

## Training a classifier with NLTK-Trainer

This can also be done as an argument in `train_classifier.py`:

```
$ python train_classifier.py movie_reviews --no-pickle --fraction 0.75 --show-most-informative 5 loading movie_reviews 2 labels: ['neg', 'pos'] using bag of words feature extraction 1500 training feats, 500 testing feats training NaiveBayes classifier accuracy: 0.726000 neg precision: 0.952000 neg recall: 0.476000 neg f-measure: 0.634667 pos precision: 0.650667 pos recall: 0.976000 pos f-measure: 0.780800 5 most informative features
Most Informative Features          finest = True          pos : neg = 13.4 : 1.0          astounding = True
pos : neg = 11.0 : 1.0          avoids = True          pos : neg = 11.0 : 1.0
inject = True          neg : pos = 10.3 : 1.0          strongest = True          pos : neg = 10.3 : 1.0
```

231 [www.it-ebooks.info](http://www.it-ebooks.info) Text Classification The Maxent and LogisticRegression classifiers In the Training a maximum entropy classifier recipe, we covered the `MaxentClassifier` class with the GIS algorithm. Here's how to use `train_classifier.py` to do this:

```
$ python train_classifier.py movie_reviews --no-pickle --fraction 0.75 --classifier GIS --max_iter 10 --min_lldelta 0.5 loading movie_reviews 2 labels: ['neg', 'pos'] using bag of words feature extraction 1500 training feats, 500 testing feats training GIS classifier ==> Training (10 iterations) accuracy: 0.712000 neg precision: 0.964912 neg recall: 0.440000 neg f-measure: 0.604396 pos precision: 0.637306 pos recall: 0.984000 pos f-measure: 0.773585
```

If you have `scikit-learn` installed, then you can use many different `sklearn` algorithms for classification.

## Introduction

www.it-ebooks.info . Or, you might want to store frequencies and probabilities in a persistent, shared database so multiple processes can access it simultaneously.

## Distributed tagging with execnet

The gateway's `remote_exec()` method takes a single argument that can be one of the following three types: A string of code to execute remotely f The name of a pure function that will be serialized and executed remotely f The name of a pure module whose source will be executed remotely f 239  
www.it-ebooks.info Distributed Processing and Handling Large Datasets We use option three with the `remote_tag.py` module, which is defined as follows: `import pickle if __name__ == '__channelexec__': tagger = pickle.loads(channel.receive()) for sentence in channel: channel.send(tagger.tag(sentence))` A pure module is a module that is self-contained: it can only access Python modules that are available where it executes, and does not have access to any variables or states that exist wherever the gateway is initially created. `channel = next(channels) ...`

## Distributed chunking with execnet

This Tree is then pickled and sent back over the channel: `import pickle if __name__ == '__channelexec__': tagger = pickle.loads(channel.receive()) chunker = pickle.loads(channel.receive())` 243 www.it-ebooks.info . This will be very similar to the tagging in the previous recipe, but we'll be sending two objects instead of one, and we will be receiving a Tree instead of a list, which requires pickling and unpickling for serialization.

## Parallel list processing with execnet

Once we have all the expected results, we can exit the gateways and return the results: `import itertools, execnet def map(mod, args, specs=[('popen', 2)]): gateways = [] channels = [] for spec, count in specs: for i in range(count): gw = execnet.makegateway(spec) gateways.append(gw) channels.append(gw.remote_exec(mod)) cyc = itertools.cycle(channels) for i, arg in enumerate(args): channel = next(cyc) channel.send((i, arg)) mch = execnet.MultiChannel(channels) queue = mch.make_receive_queue() l = len(args) results = [None] * l # creates a list of length l, where every element is None for i in range(l): channel, (i, result) = queue.get() results[i] = result for gw in gateways: gw.exit() return results` 246  
www.it-ebooks.info . `if __name__ == '__channelexec__': for (i, arg) in channel: channel.send((i, arg * 2))` To use this module to double every element in a list, we import the `plists` module (explained in the How it works...

## Storing a frequency distribution in Redis

This RedisHashFreqDist (defined in redisprob.py) sums all the values in the hash map for the N() method:

```
from rediscollections import RedisHashMap
class RedisHashFreqDist(RedisHashMap):
    def N(self):
        return int(sum(self.values()))
    def __missing__(self, key):
        return 0
    def __getitem__(self, key):
        return int(RedisHashMap.__getitem__(self, key) or 0)
    def values(self):
        return [int(v) for v in RedisHashMap.values(self)]
    def items(self):
        return [(k, int(v)) for (k, v) in RedisHashMap.items(self)]

We can use this class just like a FreqDist.
import collections, re
white = re.compile('[\s&]+')
def encode_key(key):
    return white.sub('_', key.strip())

class RedisHashMap(collections.MutableMapping):
    def __init__(self, r, name):
        self._r = r
        self._name = encode_key(name)
    def __iter__(self):
        return self.items()
    def __len__(self):
        return self._r.hlen(self._name)
    def __contains__(self, key):
        return self._r.hexists(self._name, encode_key(key))
    def __getitem__(self, key):
        return self._r.hget(self._name, encode_key(key))
    def __setitem__(self, key, val):
        self._r.hset(self._name, encode_key(key), val)
    def __delitem__(self, key):
        self._r.hdel(self._name, encode_key(key))
    def keys(self):
        return self._r.hkeys(self._name)
    def values(self):
        return self._r.hvals(self._name)
    def items(self):
        return self._r.hgetall(self._name).items()
    def get(self, key, default=0):
        return self[key] or default
    def clear(self):
        self._r.delete(self._name)
```

250 www.it-ebooks.info .

## Storing a conditional frequency distribution in Redis

We override \_\_getitem\_\_() so we can create an instance of RedisHashFreqDist instead of a FreqDist:

```
from nltk.probability import ConditionalFreqDist
from rediscollections import encode_key

class RedisConditionalHashFreqDist(ConditionalFreqDist):
    def __init__(self, r, name, cond_samples=None):
        self._r = r
        self._name = name
        ConditionalFreqDist.__init__(self, cond_samples)
    for key in self._r.keys(encode_key('%s:*' % name)):
        condition = key.split(':')[1]
        self[condition] # calls self.__getitem__(condition)
    def __getitem__(self, condition):
        if condition not in self._fdists:
            key = '%s:%s' % (self._name, condition)
            val = RedisHashFreqDist(self._r, key)
            super(RedisConditionalHashFreqDist, self).__setitem__(condition, val)
        return super(RedisConditionalHashFreqDist, self).__getitem__(condition)
    def clear(self):
        for fdist in self.values():
            fdist.clear()

An instance of this class can be created by passing in a Redis connection and a base name. Here, we'll create an API-compatible class on top of Redis using the RedisHashFreqDist from the previous recipe.
```

## Storing an ordered dictionary in Redis

Then, it implements all the key methods that require Redis ordered set (also known as Zset) commands:

```
class RedisOrderedDict(collections.MutableMapping):
    def __init__(self, r, name):
        self._r = r
        self._name = encode_key(name)
    def __iter__(self):
        return iter(self.items())
    def __len__(self):
        return self._r.zcard(self._name)
    def __getitem__(self, key):
        return self._r.zscore(self._name, encode_key(key))
    def __setitem__(self, key, score):
        self._r.zadd(self._name, encode_key(key), score)
    def __delitem__(self, key):
        self._r.zrem(self._name, encode_key(key))
    def keys(self, start=0, end=-1):
        # we use zrevrange to get keys sorted by high value instead of by lowest
        return self._r.zrevrange(self._name, start, end)
    def values(self, start=0, end=-1):
        return [v for (k, v) in self.items(start=start, end=end)]
    def items(self, start=0, end=-1):
        return self._r.zrevrange(self._name, start, end, withscores=True)
    def get(self, key, default=0):
        return self[key] or default
    def iteritems(self):
        return iter(self)
    def clear(self):
        self._r.delete(self._name)
```

254 [www.it-ebooks.info](http://www.it-ebooks.info) Chapter 8 You can create an instance of RedisOrderedDict by passing in a Redis connection and a unique name:

```
>>> from redis import Redis
>>> from redis.collections import RedisOrderedDict
>>> r = Redis()
>>> rod = RedisOrderedDict(r, 'test')
>>> rod.get('bar')
>>> len(rod)
0
>>> rod['bar'] = 5.2
>>> rod['bar']
5.2000000000000002
>>> len(rod)
1
>>> rod.items()
[(b'bar', 5.2)]
>>> rod.clear()
```

By default, keys are returned as binary strings. `items()`: This uses the `zrevrange` command to get the scores of each key in order `items()` to return a list of 2-tuples ordered by the highest score.